

Introduction to Machine Learning and Classical Supervised Learning


John Armstrong, Astronomy & Astrophysics



What is Machine Learning?

- Machine learning is the process of using statistical techniques to give computers the ability to learn how to perform a specific task *without* being explicitly programmed
 1. We use machine learning algorithms to approximate statistical models that may be too complex to work out
 2. This is done using a data-driven model approach: more data implies better model

Data-driven Models



Data-driven Models

4	13	6	11	8
A	B	A	B	A

7
?

Consistent Hypotheses:

- odd implies B
- less than 10 implies A

Data-driven Models

4	13	6	11	8
A	B	A	B	A

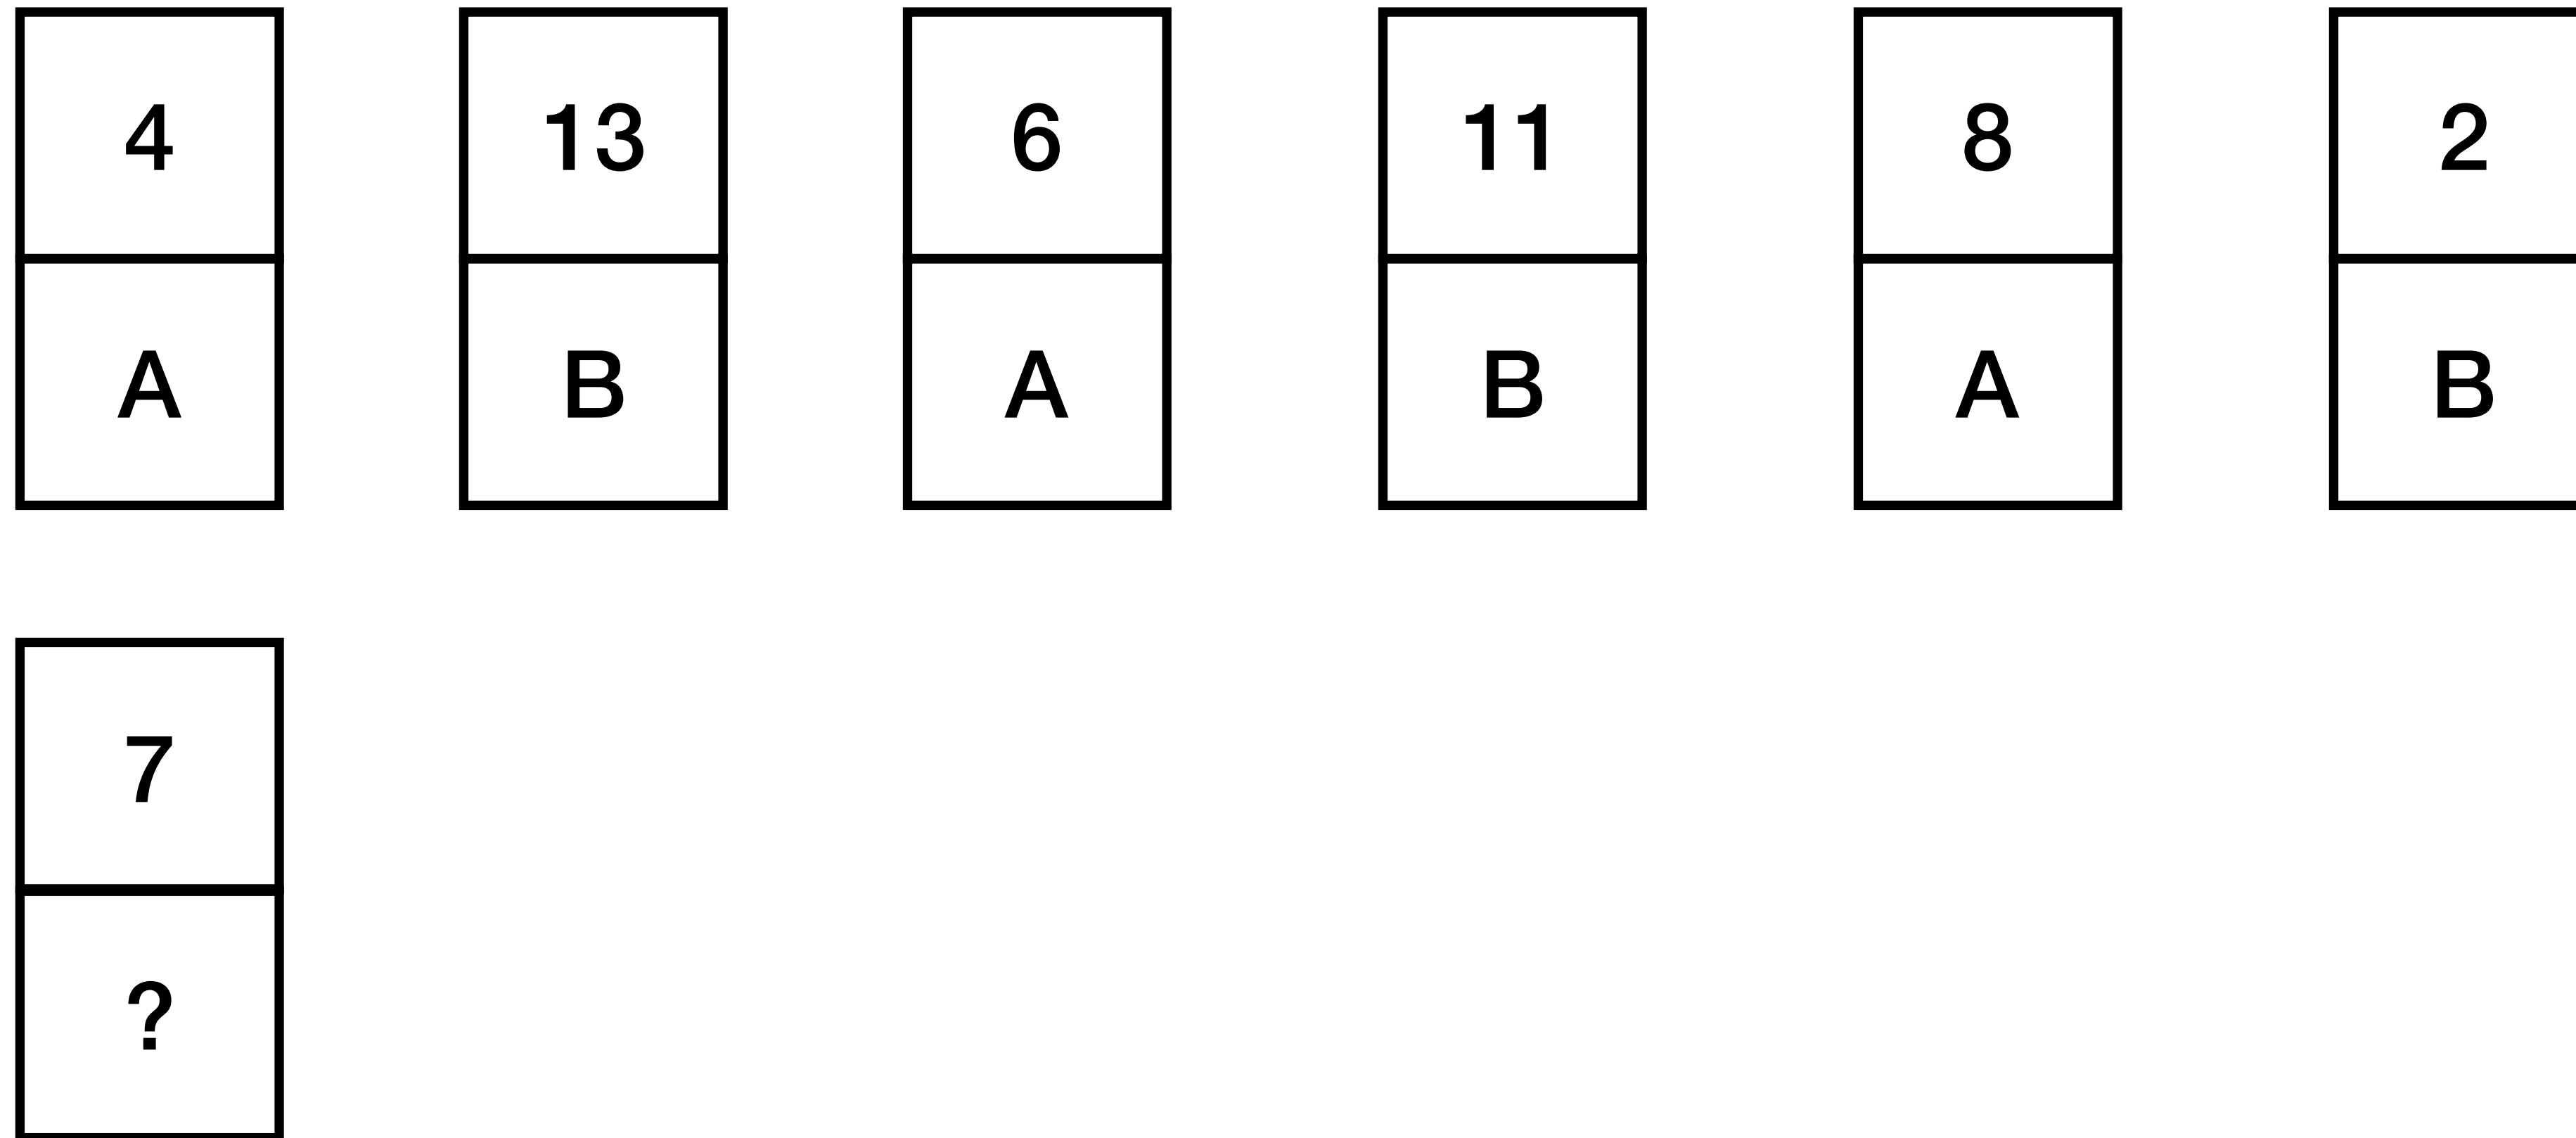
7
?

Consistent Hypotheses:

- odd implies B
- less than 10 implies A

Not enough data to get unambiguous solutions

Data-driven Models



Data-driven Models

4	13	6	11	8	2
A	B	A	B	A	B

7
?

Only consistent hypothesis:

- prime implies B

Data-driven Models

4	13	6	11	8	2
A	B	A	B	A	B

7
?

Only consistent hypothesis:

- prime implies B

Space of viable models shrinks with number of examples

- As shown, our data can highly influence the model we use to describe it
- With an exponential growth we cannot feasibly look through all of our data to look for the optimal model to use
- This is why we turn to automation and machine learning

- Given correct prior assumptions about our data and what we're looking for we should be able to automate the process of finding the most optimised model
- This can be more challenging than just applying a pre-built algorithm due to data complexity
- There are a variety of ways to get to our data driven model

Different Types of Machine Learning

	Supervised	Unsupervised
Classical	Decision Trees, SVMs, shallow neural networks	Clustering, Dimensionality Reduction
Deep	Multi-layer ANNs/ CNNs, ResNets	GANs, VAEs, INNs

Different Types of Machine Learning

	Supervised	Unsupervised
Classical	Decision Trees, SVMs, shallow neural networks	Clustering, Dimensionality Reduction
Deep	Multi-layer ANNs/ CNNs, ResNets	GANs, VAEs, INNs

Supervised Learning I

- This is when you know the structure of your data and want to create a model that will predict a characteristic of this data for an unknown point
- As such, these methods are good for coping with two types of problems: classification and regression
- Understanding whether your machine learning task is classification or regression is the key for selecting the right algorithms to use.

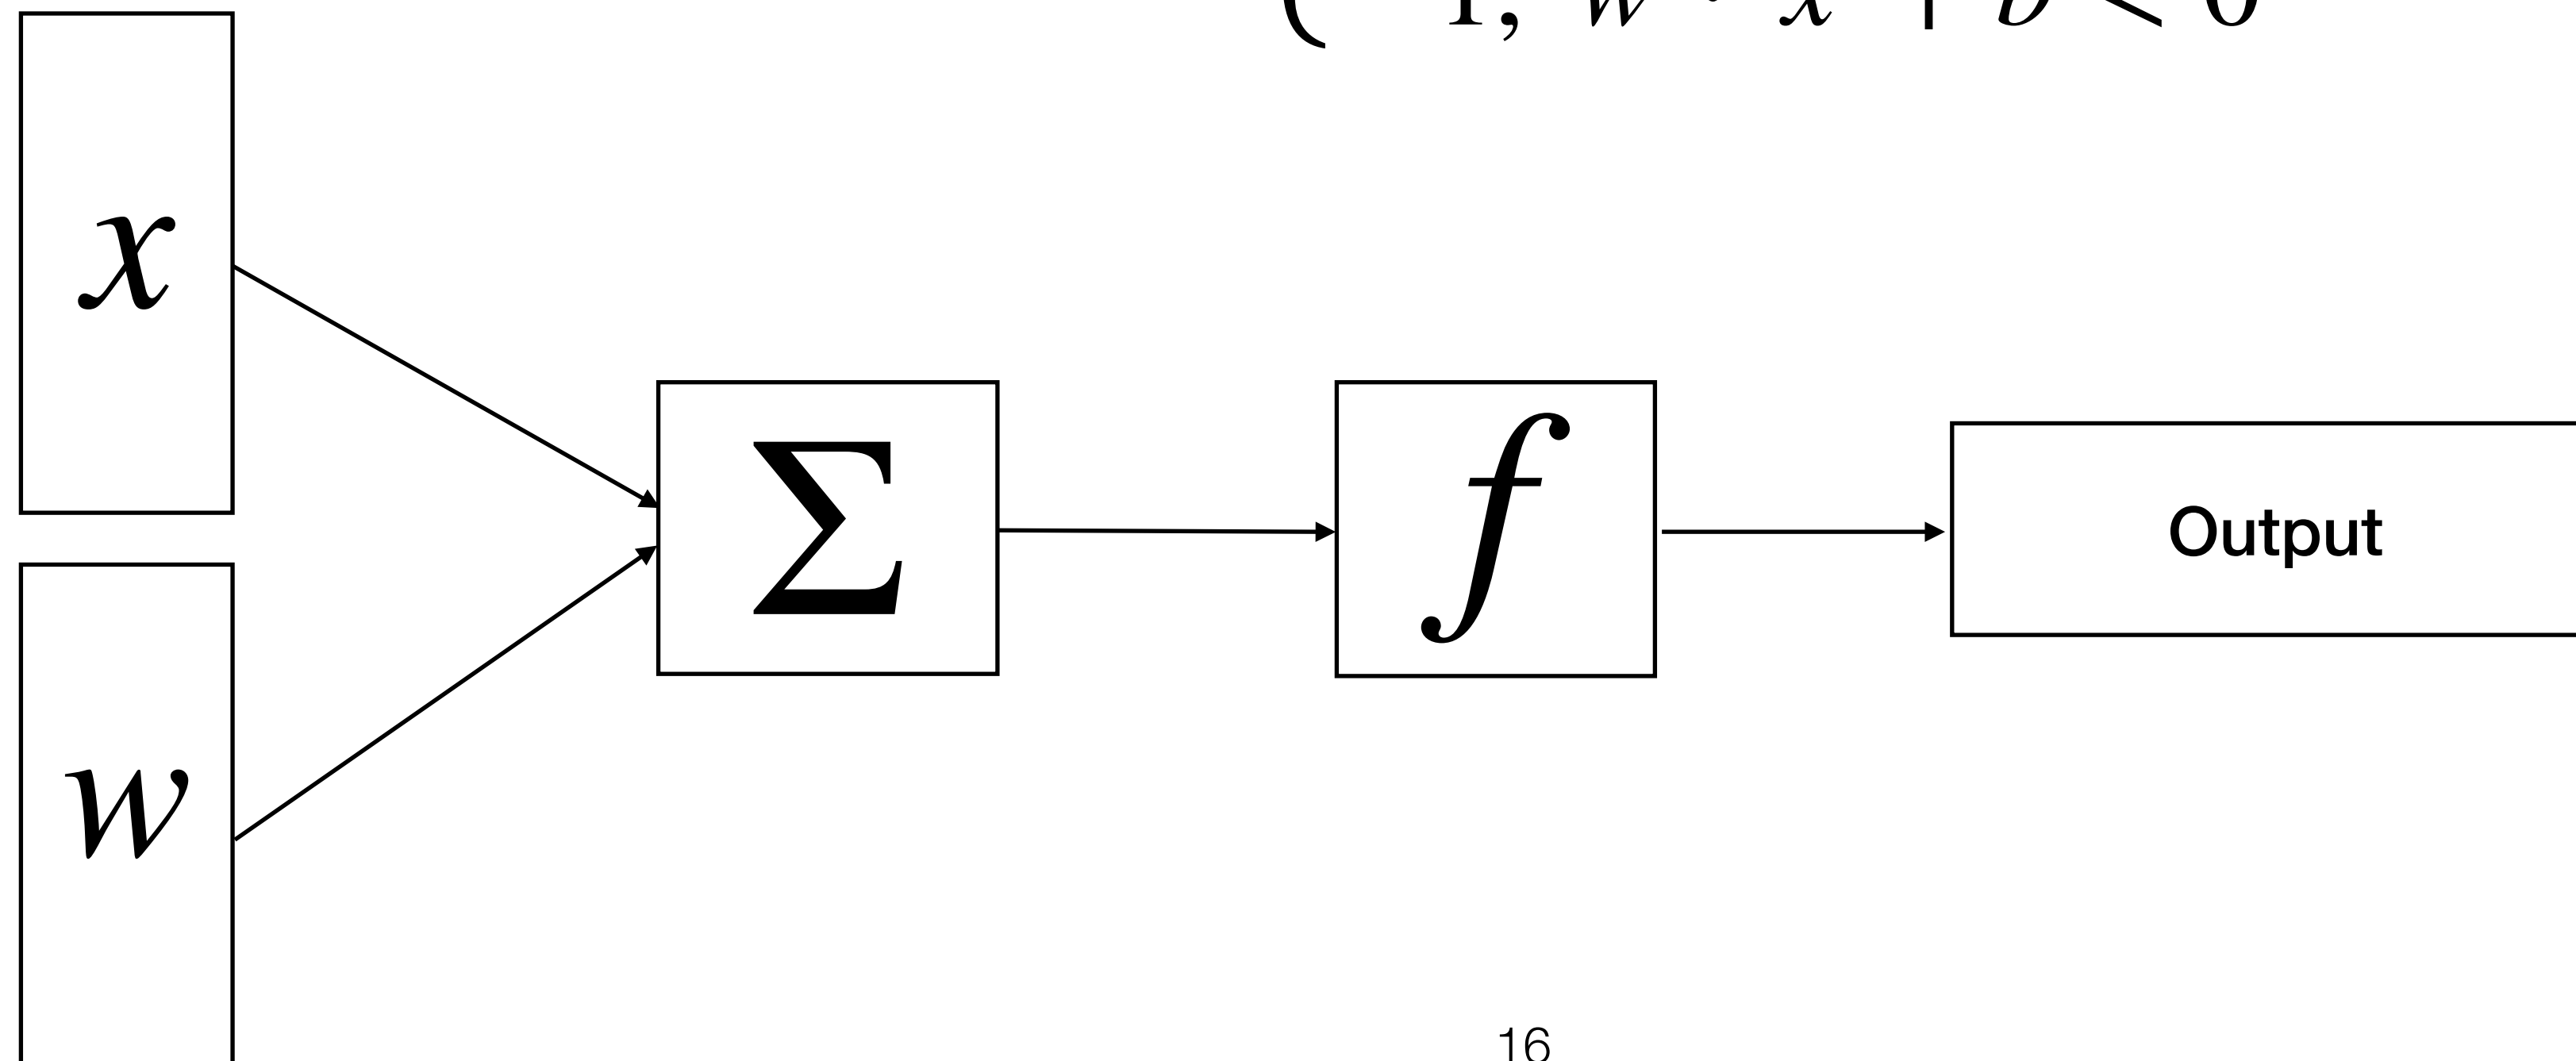
Supervised Learning II

- Typical setup for a supervised learning problem:
 - large dataset with known labels split between a training and validation set (90%/10%)
 - some kind of classification or regression algorithm
 - a non-ambiguous problem to solve
- Next we will talk about different algorithms for supervised learning



I. The Perceptron

- Based on biological model of a neuron (at the time of conception)
- Binary classification/regression only. Unable to solve non-linear problems.
- Mathematically simple: $y = \begin{cases} 1, & \vec{w} \cdot \vec{x} + b \geq 0 \\ -1, & \vec{w} \cdot \vec{x} + b < 0 \end{cases}$



- Trained with backpropagation and Stochastic Gradient Descent (SGD).

$$w_i(t) = w_i(t - 1) + \eta \nabla_w L(w; x)$$

New, updated weight for
current timestep

Weight from previous
timestep

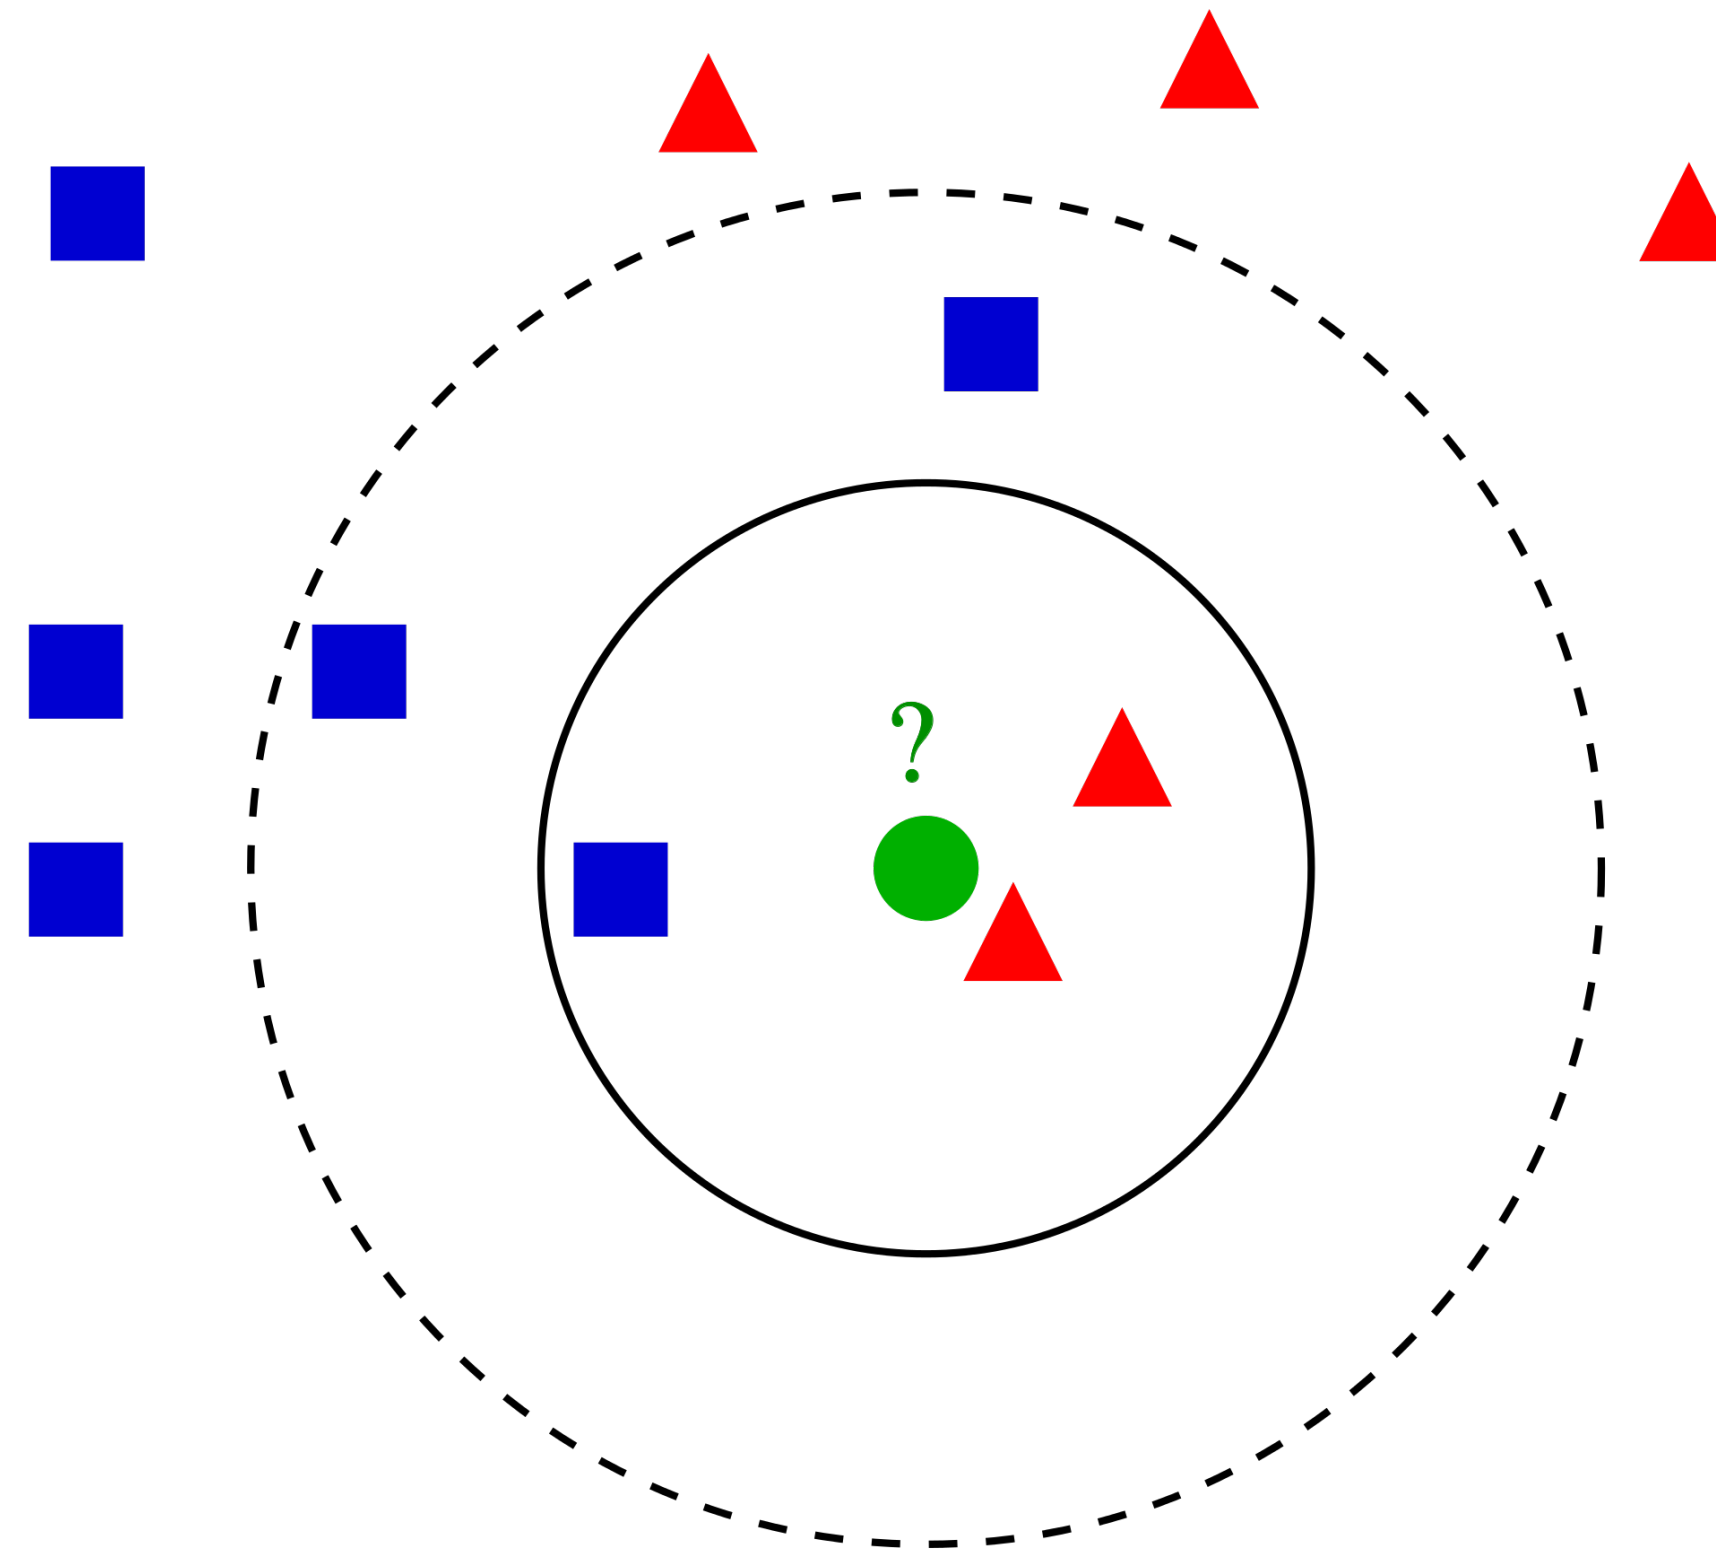
Learning Rate

Loss function

II. K-Nearest Neighbours (K-NN)

- For a point in your feature space, take the k nearest points to it and calculate a distance metric.
- In classification, a majority vote decides on the class label of the point
- In regression, an average of the distances decides where the point lies
- This implicitly calculates the decision boundary in your data

- k-NN is known as instance-based or lazy learning as it only ever approximates the function locally and it has no need to compute it globally



III. Gaussian Processes

Gaussian Processes I

- Each class in the dataset is distributed normally
- The prior of the process is then the joint distribution of the the classes in the problem i.e. our prior is a multivariate Gaussian
- We then want to use the covariance matrix between the points (often referred to as the kernel) to constrain the posterior distribution of our data

- More generally, we assume a Gaussian prior of a set of functions which may represent our problem and use the covariance of the data to constrain the posterior on this set of functions

$$p(y^* | x^*, f(x), x) = N(y^* | A, B)$$

$$A = K(\theta, x^*, x) K(\theta, x, x')^{-1} f(x)$$

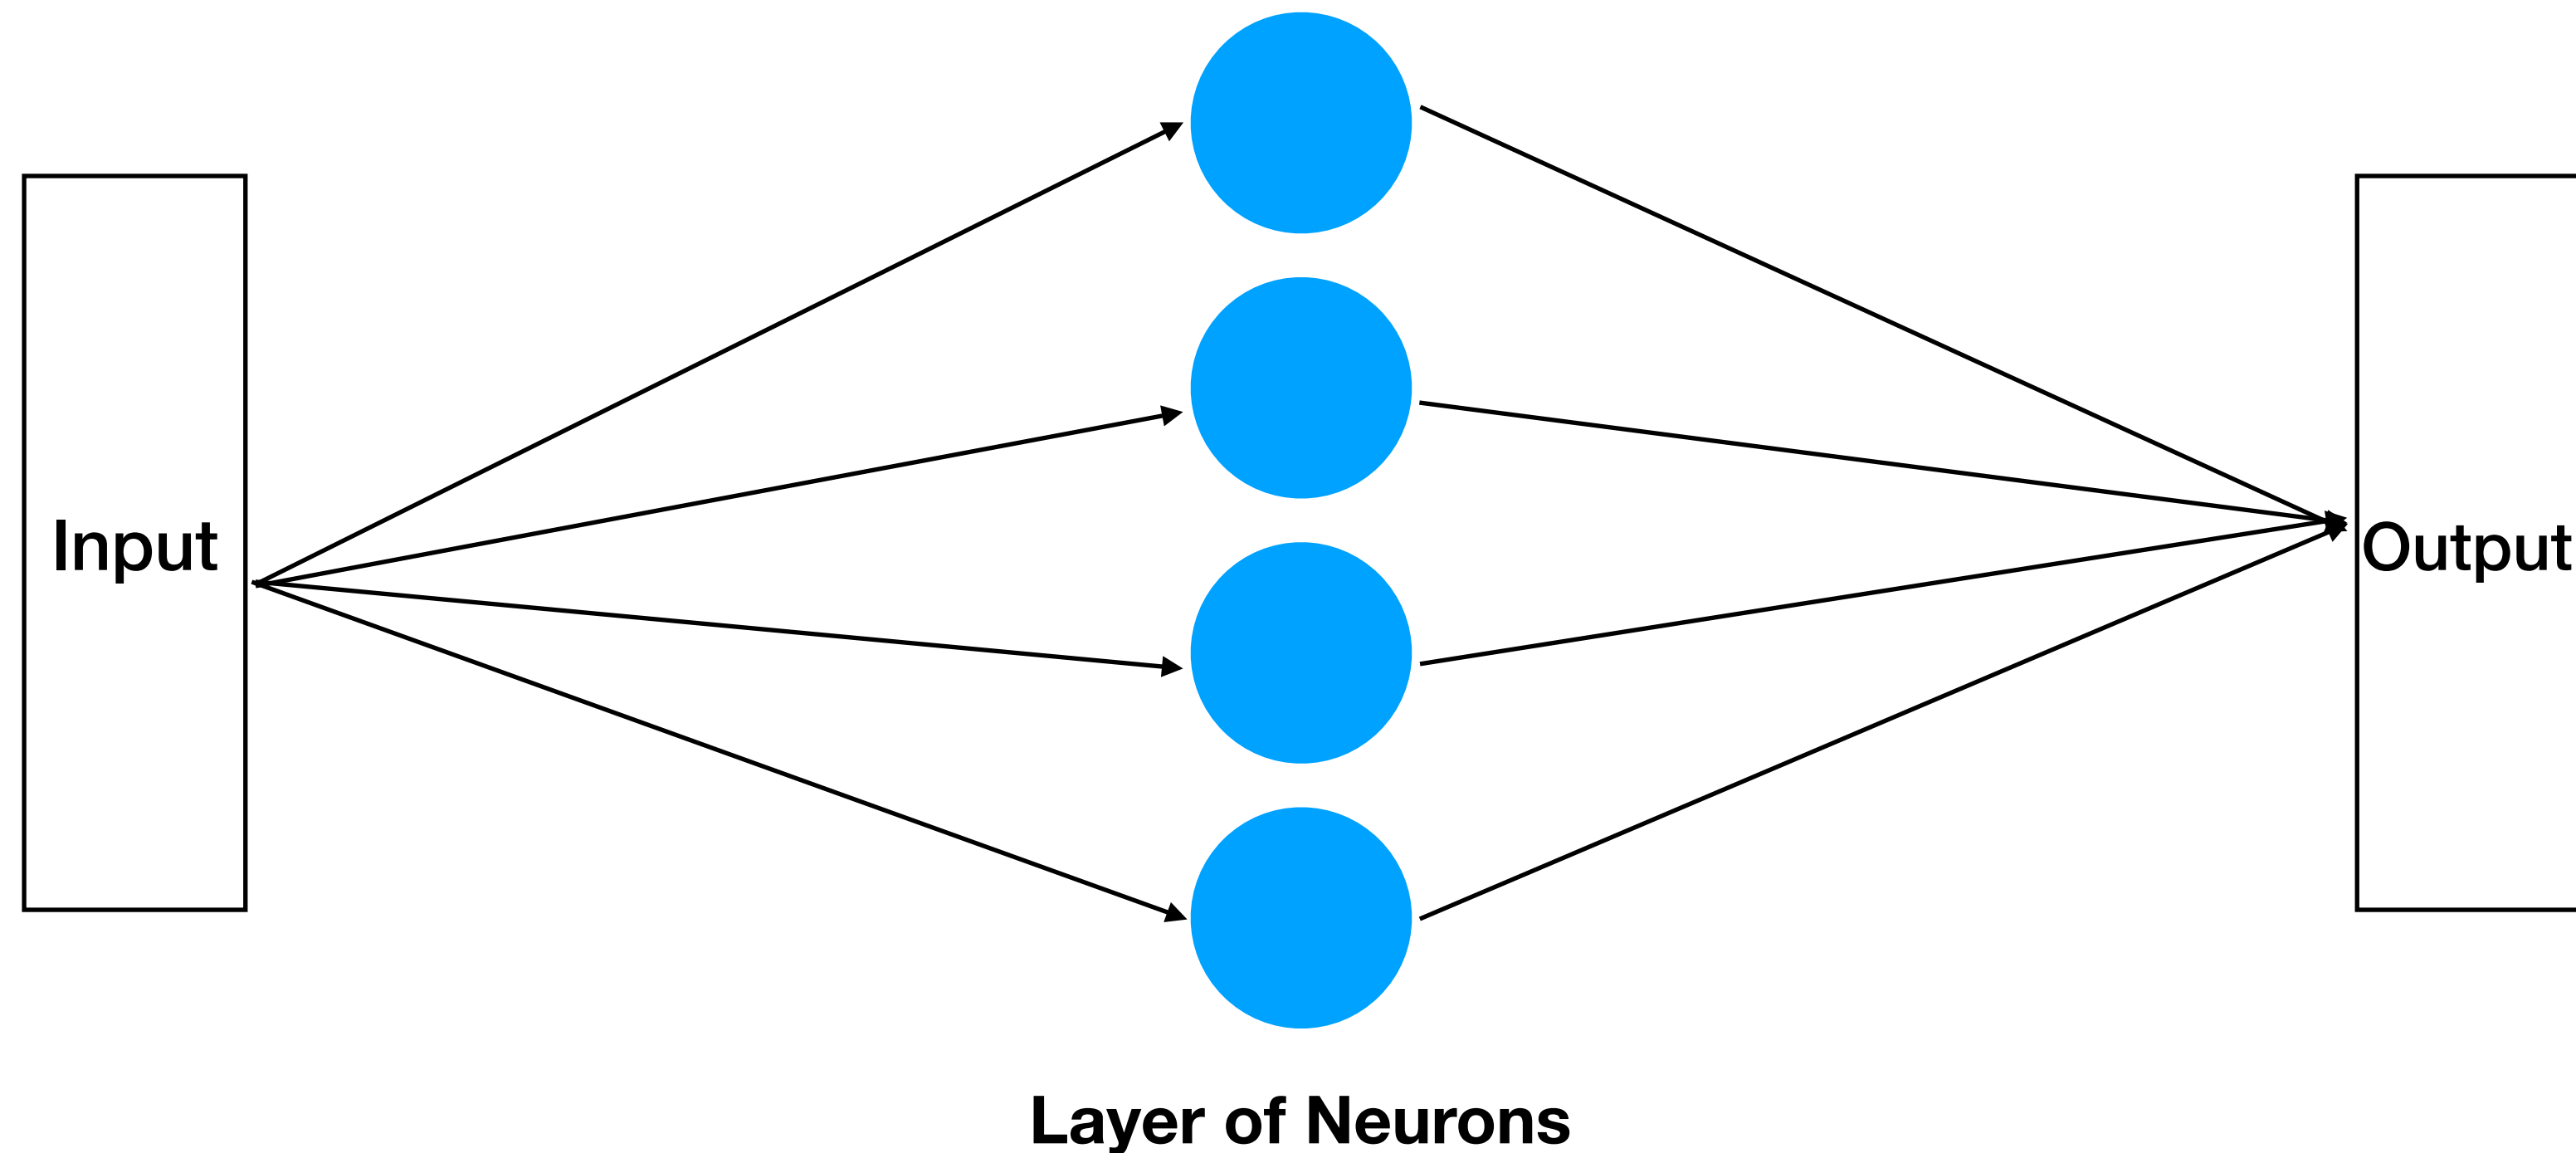
$$B = K(\theta, x^*, x^*) - K(\theta, x^*, x) K(\theta, x, x')^{-1} K(\theta, x^*, x)^T$$

IV. Support Vector Machines (SVM)

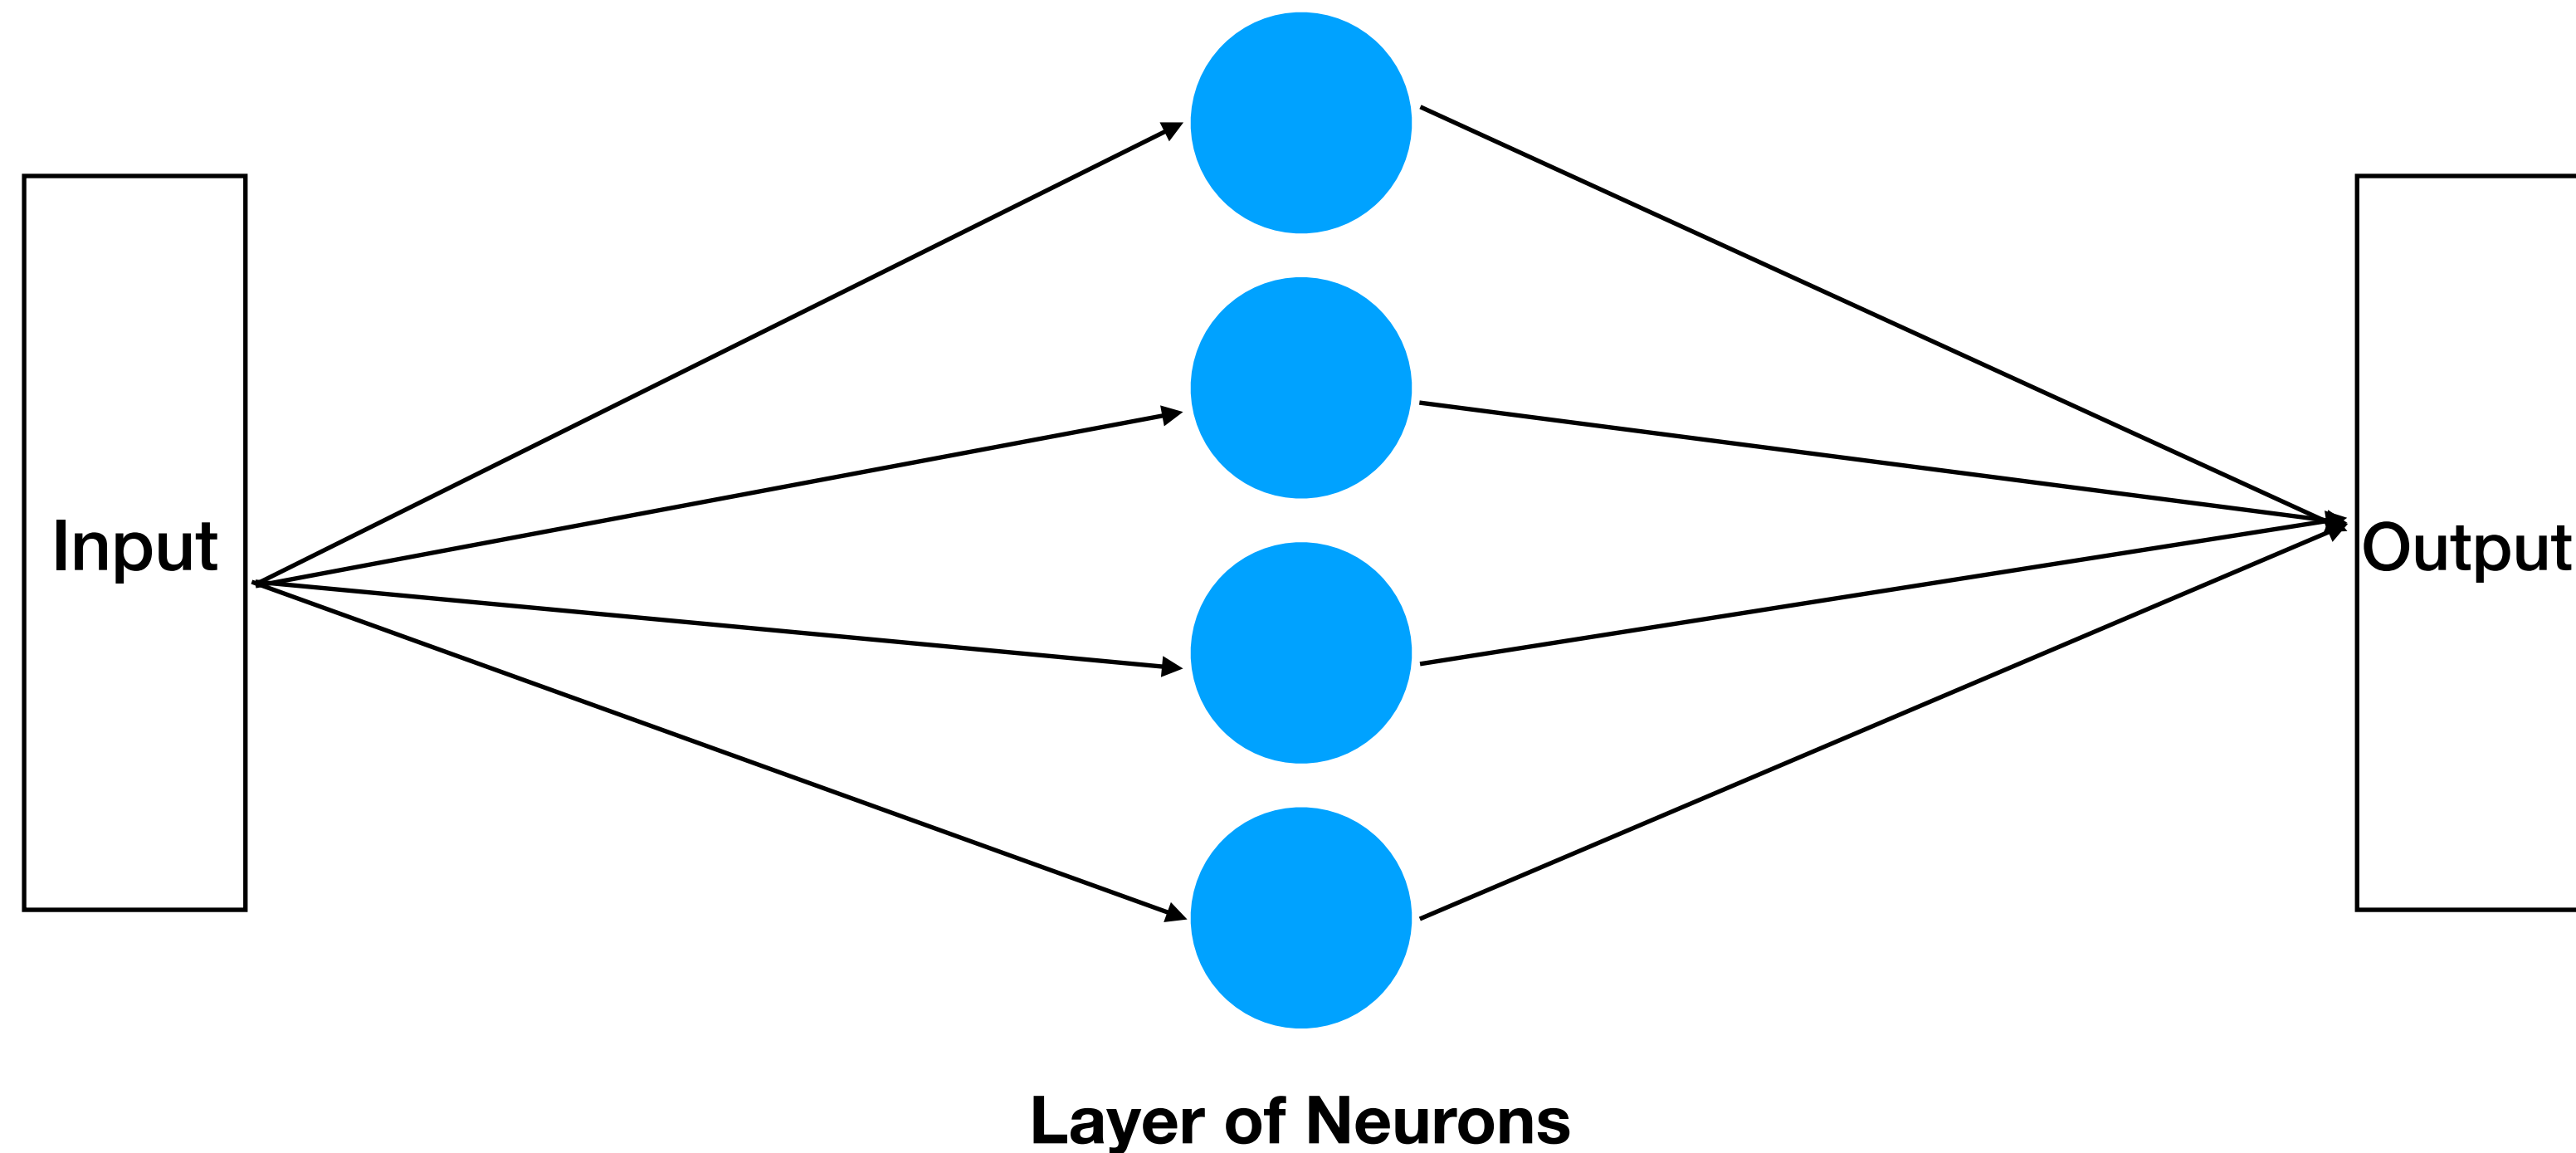
- Attempts to find a hyperplane which separates the classes in our dataset
- This is done via the use of “support vectors”
- For non-linear classification/regression tasks, SVMs find a non-linear transformation into hyperspace such that a linear decision boundary between the sets of data can be defined

V. Neural Networks (NNs)

- We will be specifically talking about shallow (single-layer) neural networks (for deep neural networks stick around until the last session)
- These can be good for approximating functions and classification



- To train a neural network, we need to define a loss function and an optimiser (often SGD).
- NNs are difficult to train but can give incredible results when trained properly.



- Clone the repository at <https://bit.ly/2D04zE1> if you haven't already
- Follow instructions to set up Python virtual environment
- We will be using 2 machine learning libraries: scikit-learn and PyTorch
- scikit-learn has most of the classical machine learning techniques built-in so no need to try and code your own
- PyTorch is a neural network equivalent of numpy and conversion between the two is easy: numpy array to Torch tensor is `y = torch.from_numpy(x)` and Torch tensor to numpy array is simply `y = x.numpy()`.